

Week 1: Introduction to Machine Learning

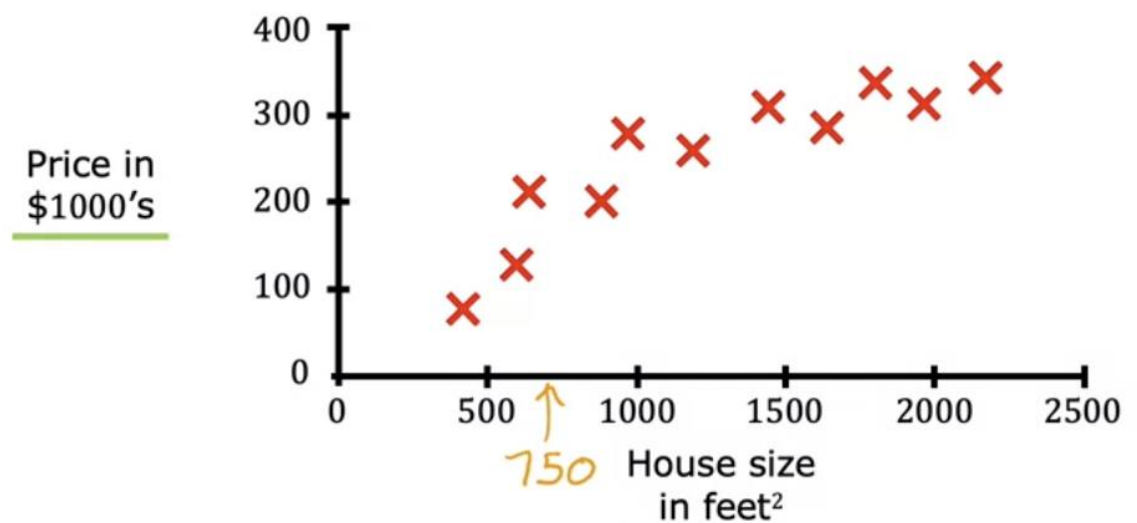
Supervised vs. Unsupervised Machine Learning

→ Supervised Learning

◆ Algorithms learn **input → output** or **X (input) → Y (output label)**

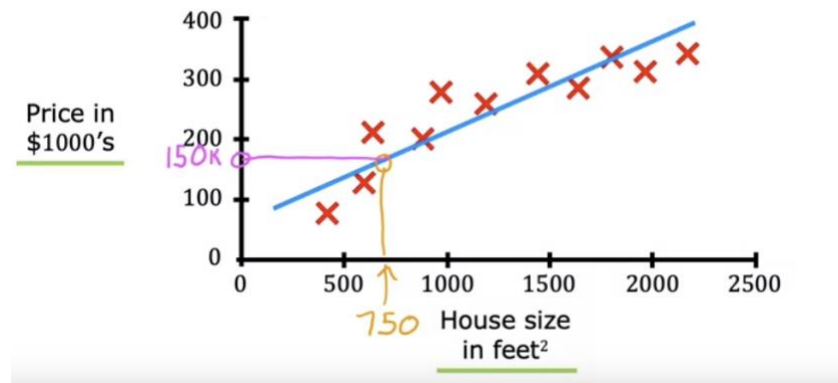
- You give the algorithm examples with the correct output so the algorithm can learn for a given input
 - → over time the algorithm will be able to take the input variable (X) without the output variable and give a reasonable output due to the learning it did
 - Algorithms after learning can take a new input variable (X) and produce a new output variable (Y) that is reasonable
- Example: Ad, User info (X) → Click? (0/1) – *Application is online advertising*
 - Train algorithm with information regarding user info and the ad and whether or not they click the advertisement. If they do, you know what users are likely to click on a certain kind of advertisement

◆ Example

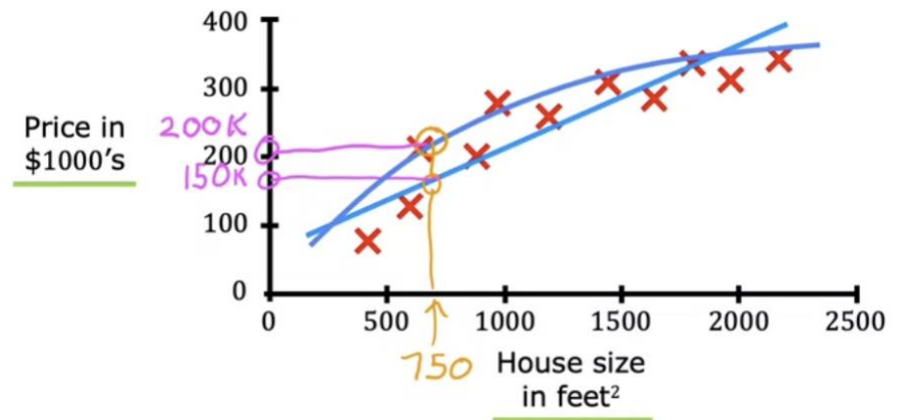


- There are some input variables and output prices already listed
- A friend wants to know the price for a 750 ft² home

- Algorithm can provide a straight line fit to determine the price of the house = approximately \$150K



- Can also use a curved line - Price of house comes out closer to \$200 K

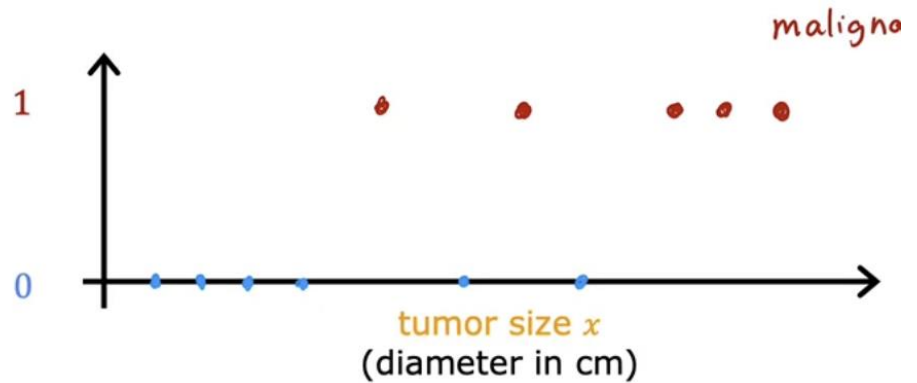


- This is supervised learning because we gave the algorithm a data set with answers to predict a new X value
 - In this example we used a process called **regression** in order to get an output of infinitely many possible values
 - ◆ *One of the major supervised learning algorithms*
- ◆ Classification algorithm
 - Example: Trying to classify whether tumors are benign or malignant

- Given data set to algorithm

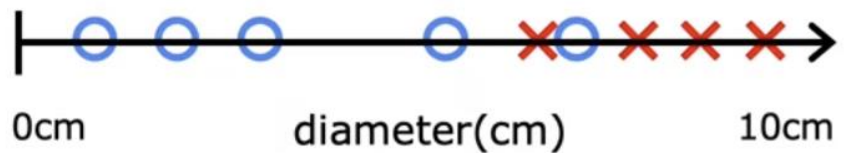
size	diagnosis
2	0
5	1
1	0
7	1
⋮	

- ◆ Size of Tumor vs Diagnosis (0 = benign, 1 = malignant)



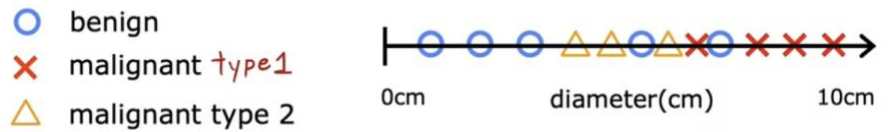
-

- ◆ Blue dots on 0 axis = benign, Red dots on 1 axis = malignant
- ◆ In this algorithm it is different since the algorithm can only predict to classify two possible diagnoses/output variables. In regression there are an infinite number of output variable possibilities
- ◆ Another way to show this is:

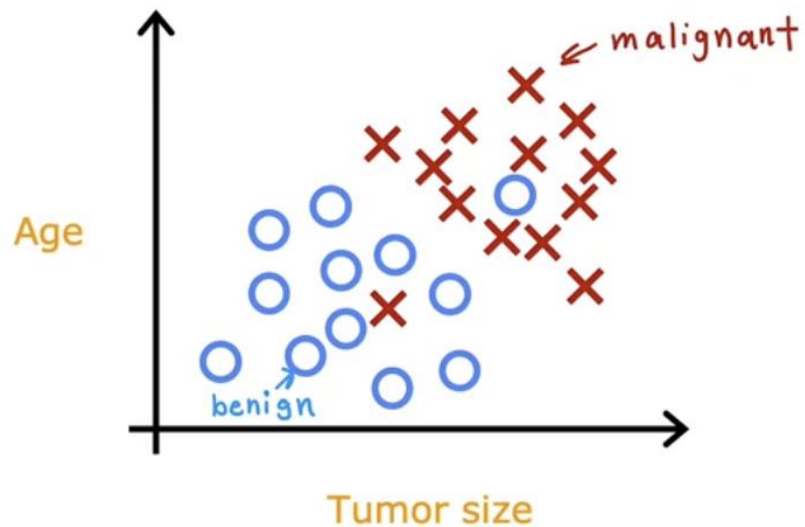


- Algorithm will try to predict whether a new tumor (not in the data set is benign or malignant)

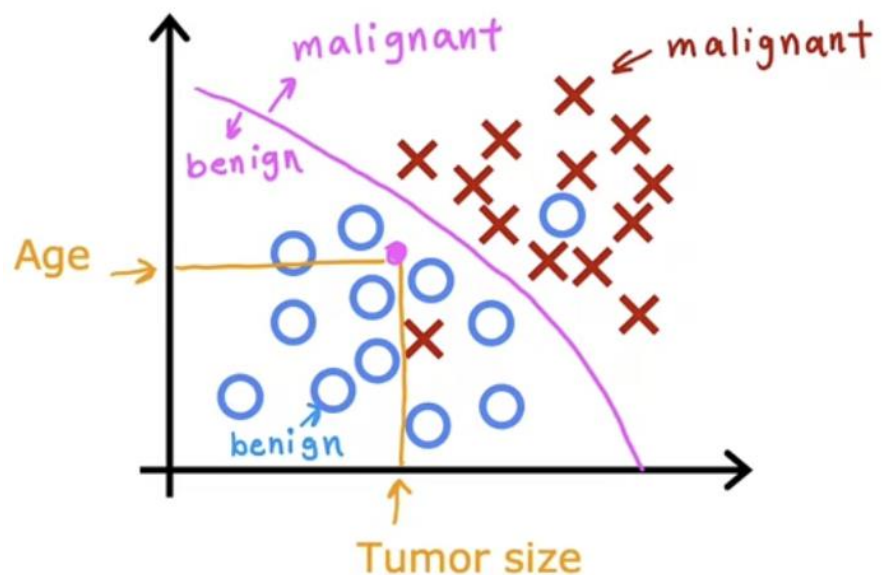
- ◆ Classification can also provide an output where there is more than two options



- In classification, output is often referred to as class or category
 - ◆ Could be non-numeric
- Example: **Two or more inputs** could also be used to predict whether a tumor is malignant or benign
 - Age and tumor size are inputs

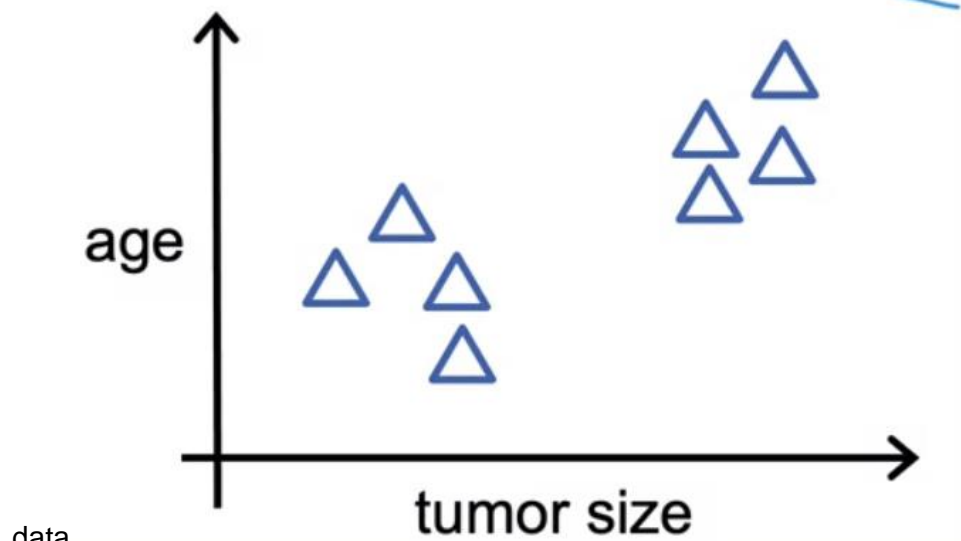


- To predict whether a person has a tumor or not, the learning algorithm might use a boundary approach

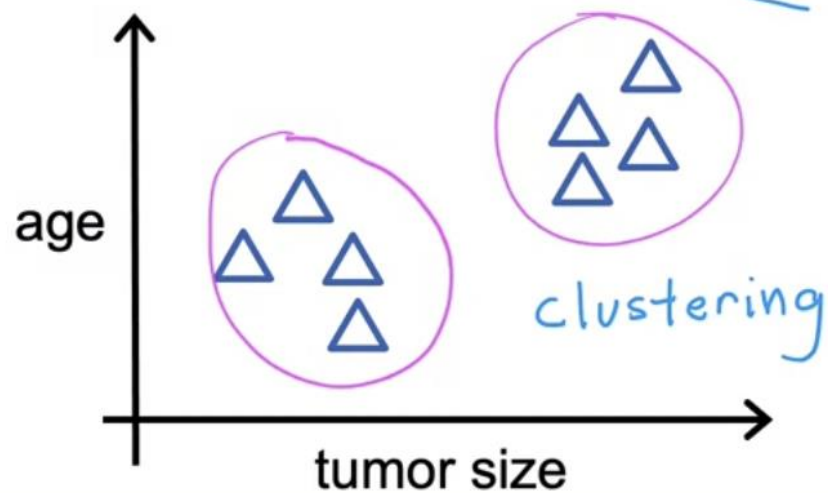


→ Unsupervised Learning

- ◆ Data is given without any output labels (Y value), but is given input values (X values)
 - Example: Not given information regarding whether tumors are benign or malignant. Our goal is just to find something interesting in the unlabeled

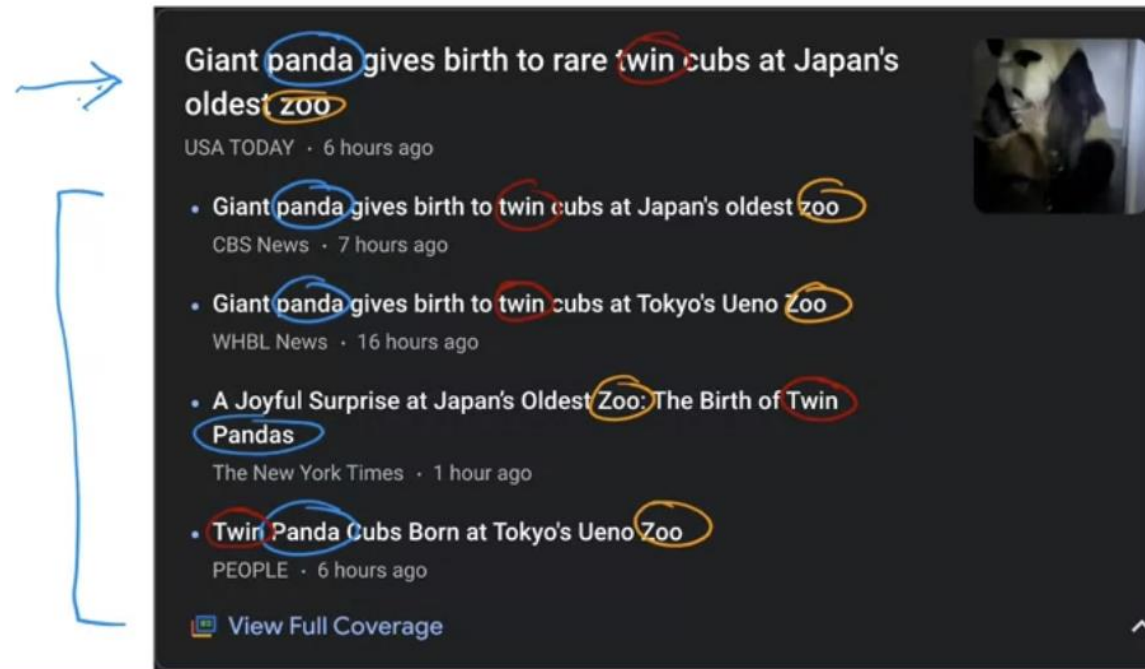


- Algorithm may decide data can be assigned to different clusters



- ◆ This is known as a **clustering algorithm** - grouping related values together

- Example: Google news using common words to cluster articles together



- ◆ Another possible algorithm - **Anomaly Detection**
 - Used to find usual data points
- ◆ Another possible algorithm - **Dimensionality Reduction**
 - Take big data set → compress data set using fewer numbers (without losing information)

Practice Quiz: Supervised vs. unsupervised learning

1. Which are the two common types of supervised learning? (Choose two)

☐ Clustering

☒ Classification

☒ Regression

2.

Which of these is a type of unsupervised learning?

☐ Classification

☒ Clustering

☐ Regression

Regression Model

→ Linear Regression Model

- ◆ Can fit a model to a straight line (Example below)



- Cost of a 1250 ft² house is ~220k
- Example of a supervised learning model since the data has “right answers”
 - Regression model since it predicts a number
 - ◆ Infinitely many possible outputs
- ◆ Could be helpful to see the data as a data table

Data table

size in feet ²	price in \$1000's
2104	400
1416	232
1534	315
852	178
...	...
3210	870

in

◆ Terminology and Notation

set:	x	y
→	size in feet ²	→ price in \$1000's
(1)	2104	400
(2)	1416	232
(3)	1534	315
(4)	852	178
...
(47)	3210	870

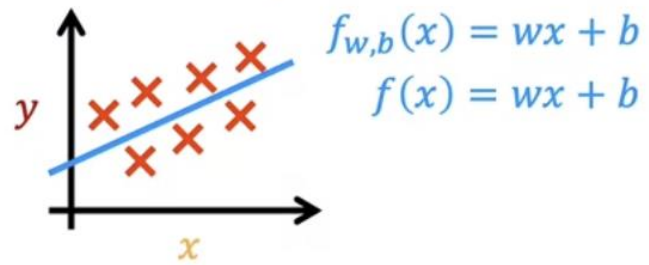
$m = 47$

$x^{(1)} = 2104$ $y^{(1)} = 400$

$(x^{(1)}, y^{(1)}) = (2104, 400)$

- Terminology
 - The data set that you use to train a particular model is called a **training set**
- Notation
 - x = “input” variable or feature
 - y = “output” variable or “target variable”
 - m = number of training examples
 - (x, y) = single training example
 - $(x^{(i)}, y^{(i)})$ = i^{th} training example
 - ◆ i refers to a specific row in the table or the i^{th} training example
 - ◆ Index of the training set
 - ◆ **Not exponent**
- ◆ Supervised learning has both input and outputs in a *training set* → produce a *training algorithm* → produces a *function* (f) or hypothesis or model
 - The *function* takes a new input (x) and produces a new output prediction (\hat{y})
 - The output prediction is the estimated value of y (target)
 - ◆ This is the actual value of y
 - How to represent f ?
 - To represent as a straight line
 - ◆ $f_{w,b}(x) = wx + b$ or could be written as $f(x)$
 - This will give us the prediction of \hat{y}

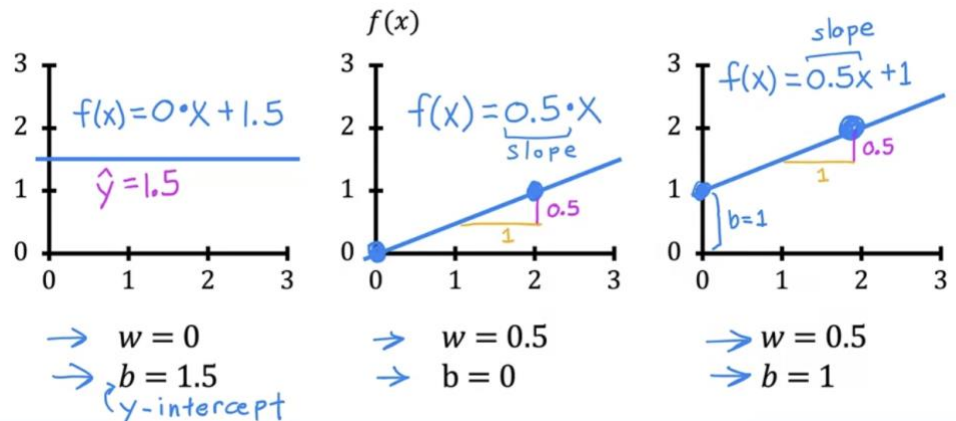
- Generates a best fit line



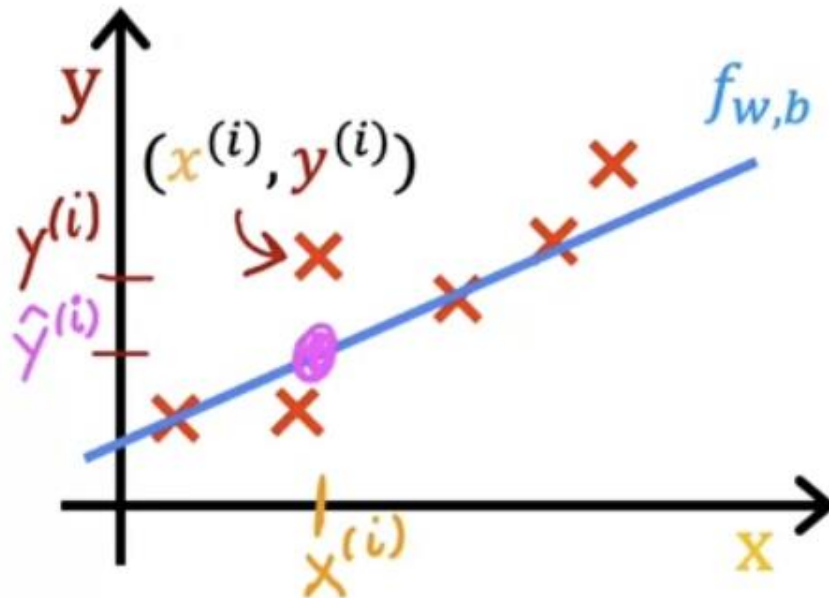
- Linear regression with one variable (single feature x) – **univariate linear regression**

→ Cost Function Formula

- ◆ Going back to the linear regression function model ($f_{w,b}(x) = wx + b$)
 - w, b are parameters/coefficients/weights
 - More specifically w is known as the **slope** and the b value is the **y-intercept**
 - Variables you can adjust during training to improve the model
 - Examples:



- Example with training set:



- The line passing through the points in the data set is roughly passing through the training examples
- To predict a \hat{y} value: $\hat{y}^{(i)} = f_{w,b}(x^{(i)})$ or $f_{w,b}(x^{(i)}) = wx^{(i)} + b$
- To find the best fit line \hat{y} should be close to $y^{(i)}$ for all $(x^{(i)}, y^{(i)})$

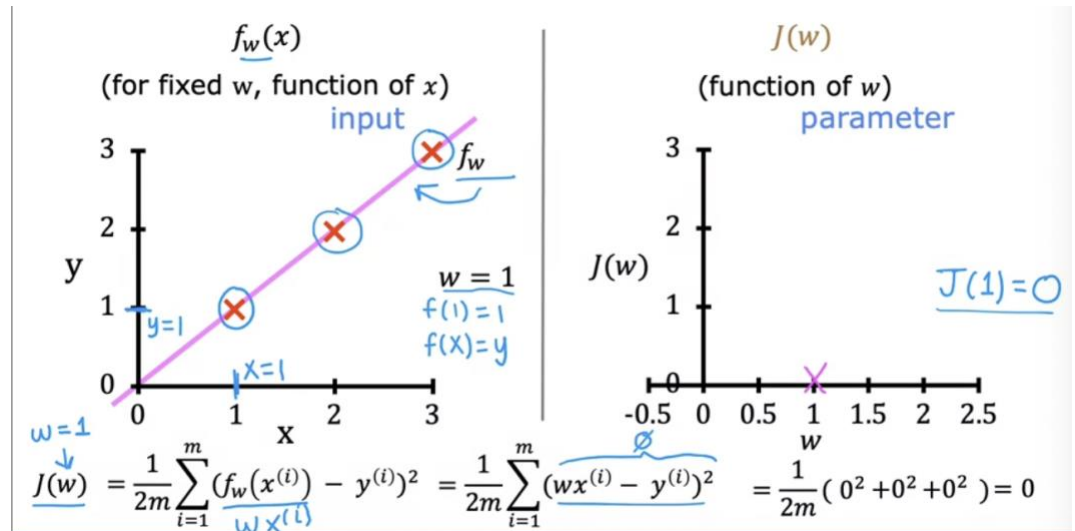
◆ Cost function

- $\hat{y} - y = \text{error}$ error \rightarrow how far off from the target the predicted value is
- $J(w, b) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2$ or $J(w, b) = \frac{1}{2m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})^2$
 - m = number of training examples
 - To prevent the summation function from getting bigger we divide it by the number of training examples (m)
 - ◆ In the function above you can see the summation is divided by $2m \rightarrow$ this is because to make machine learning a little easier

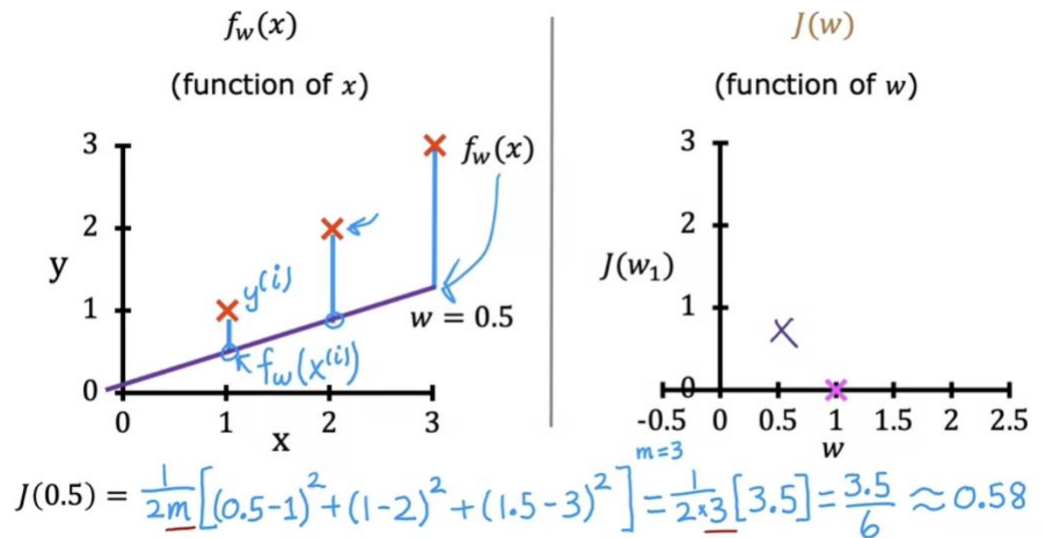
\rightarrow Cost function intuition

- ◆ The goal of is to minimize $J(w, b)$ by adjusting the values of w and b
- ◆ Example: $f_w(x) = wx$ (in this case $b = 0$)
 - \rightarrow Cost Function: $J(w) = \frac{1}{2m} \sum_{i=1}^m (f_w(x^{(i)}) - y^{(i)})^2$
 - $f_w(x^{(i)}) = wx^{(i)}$
 - Goal: Find value of w that minimizes $J(w)$

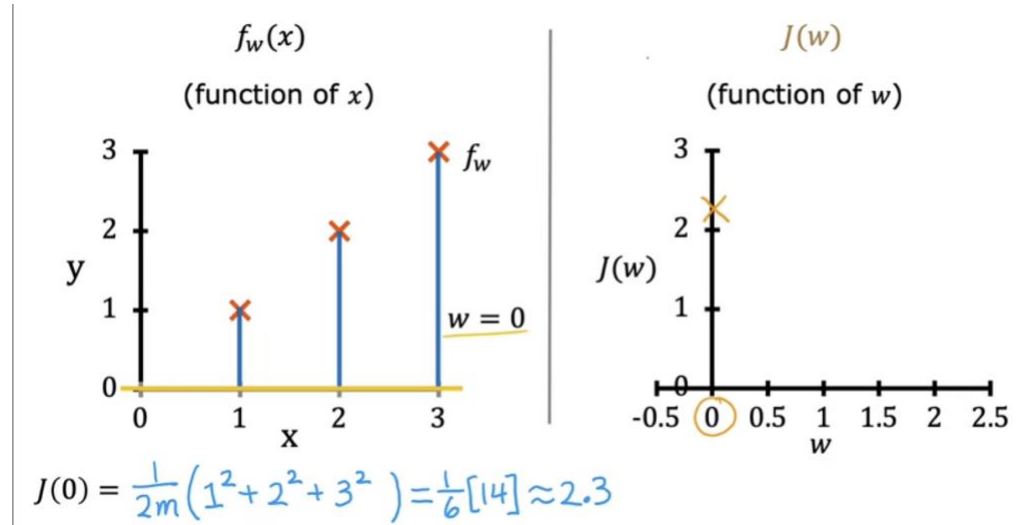
- Example: For $w = 1$



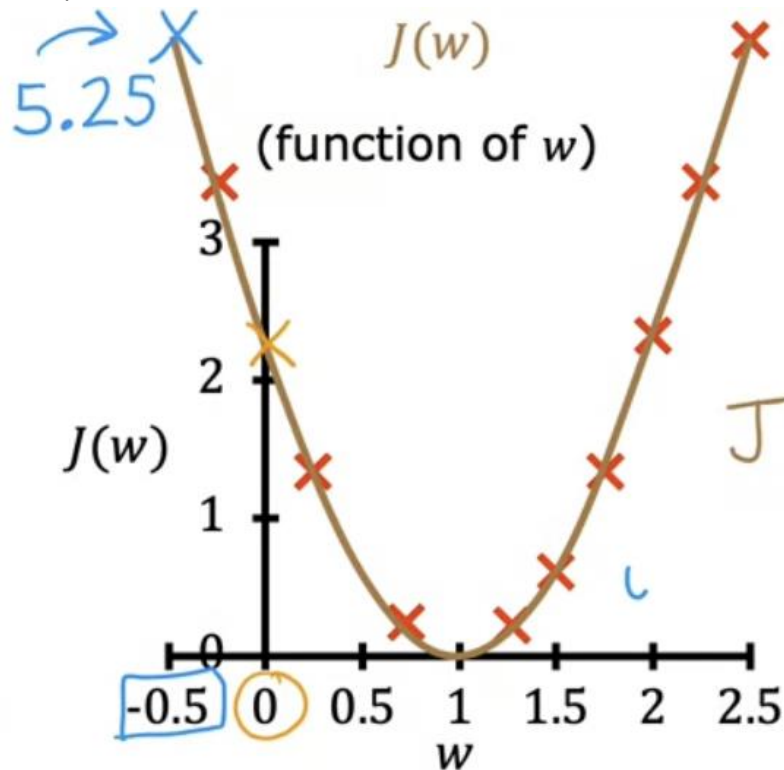
- Example: For $w = 0.5$



- Example: $w = 0$



- Example: Other w values



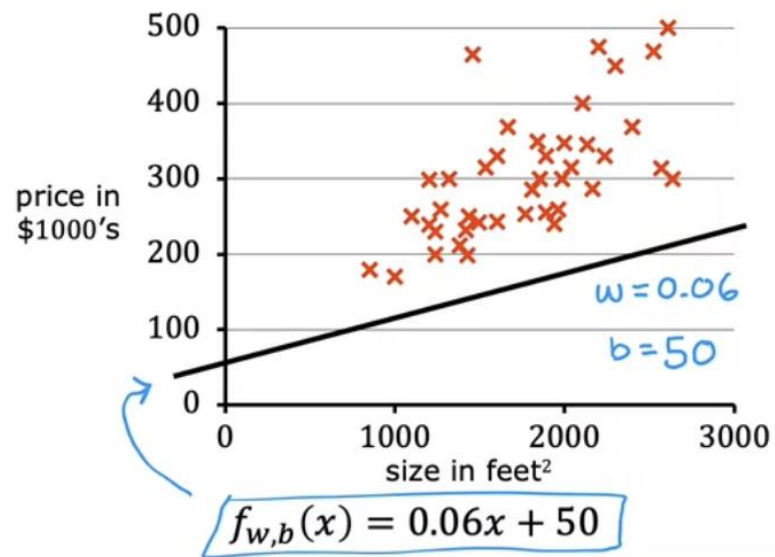
- ◆ How to choose w ?

- Choose w to minimize $J(w)$ to minimize the square errors
 - For this example you would choose $w = 1$

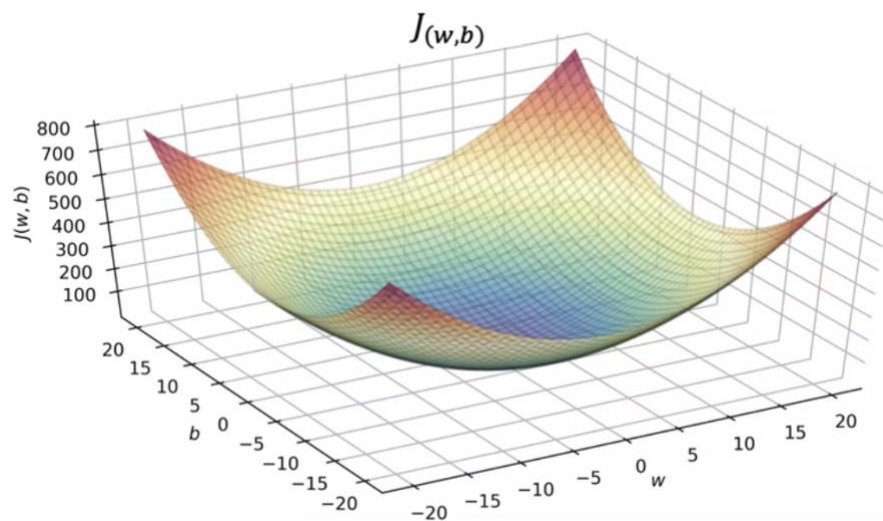
→ Visualizing the Cost Function

- ◆ We are going to now module w and b

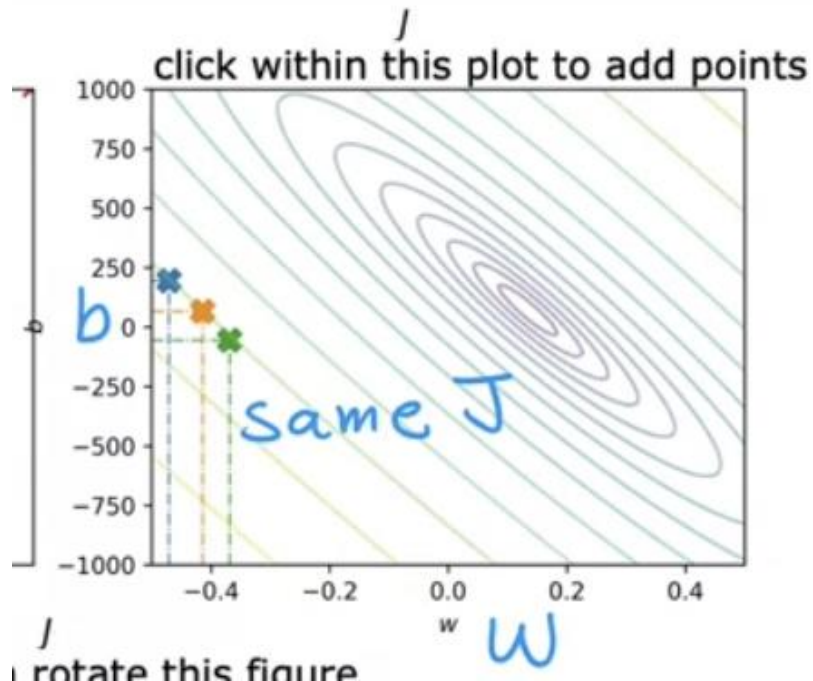
- ◆ Example: Using this data set and setting $w = 0.06$ and $b = 50$ we get the following line



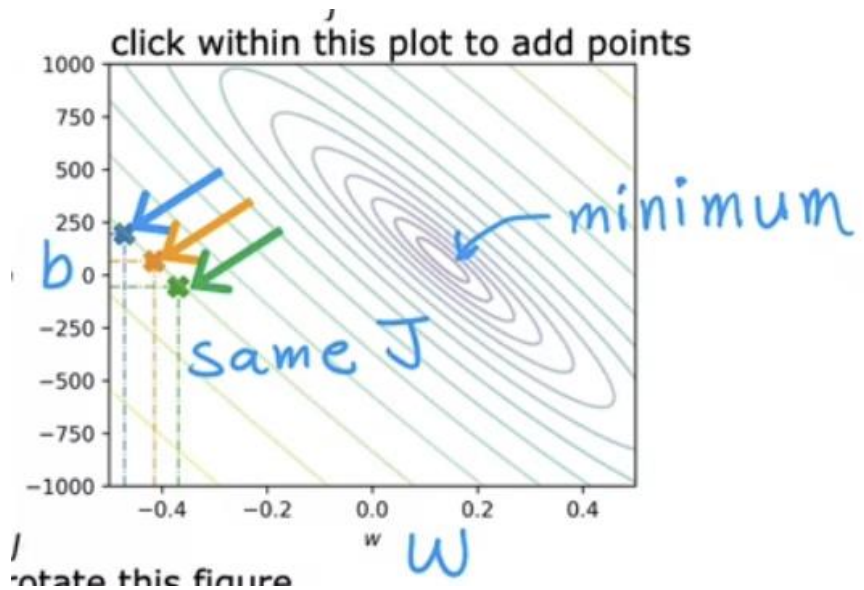
- → We get this 3-dimensional graph (since we have to now plot w and b) for $J(w, b)$



- While 3-D plots are nice, we can also use a contour plot to visualize



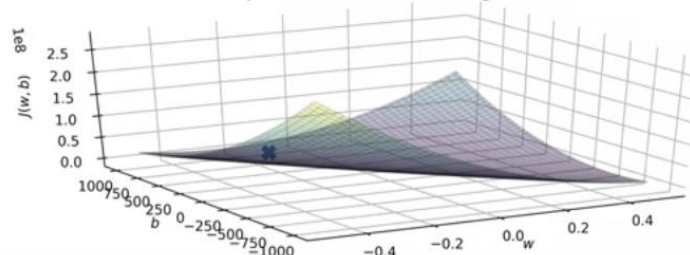
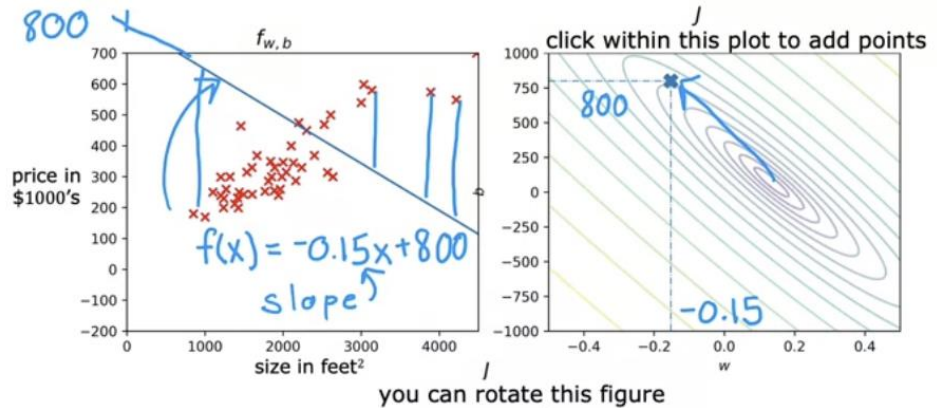
- Y axis is b and X axis is w
- Ovals show center points on 3D surface which are the same height or the same cost function (J)
 - ◆
- To get contour plot you have a 3-D plot and slice contours horizontally
 - ◆ This way you get all values with the same height
- Bottom of the bowl where J is at a minimum, indicated ideal w and b values



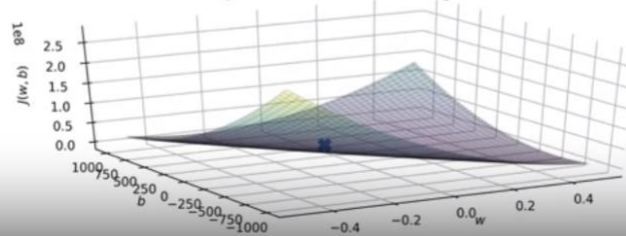
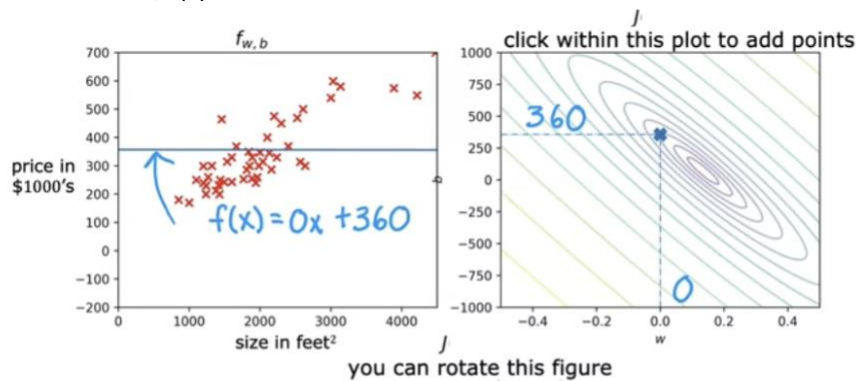
→ Visualization Examples

- ◆ Using the function, $f(x) = -0.15x + 800$, we can see there is a high error margin since the data points are far from the function we chose

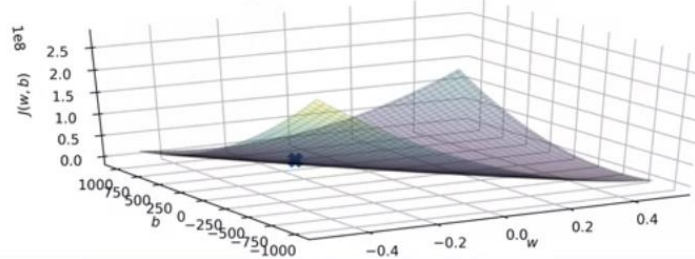
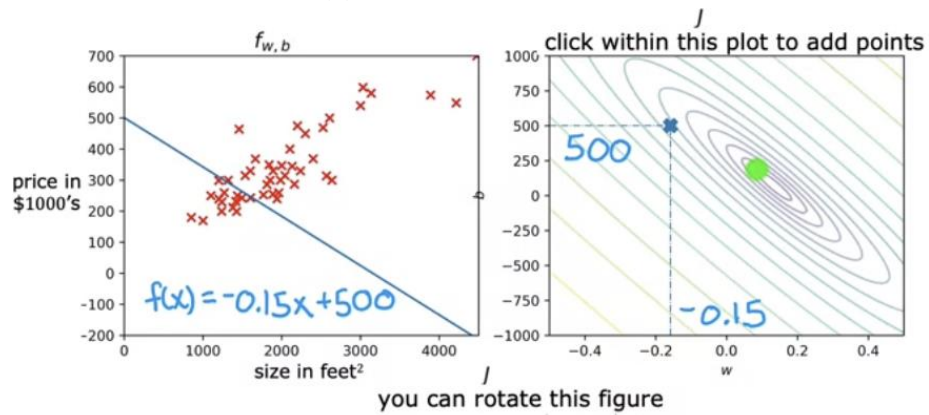
- $w = -0.15$
- $b = 800$
- → Get this contour plot



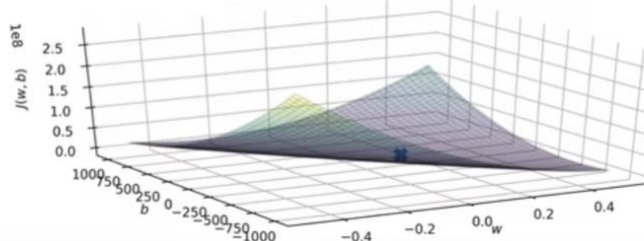
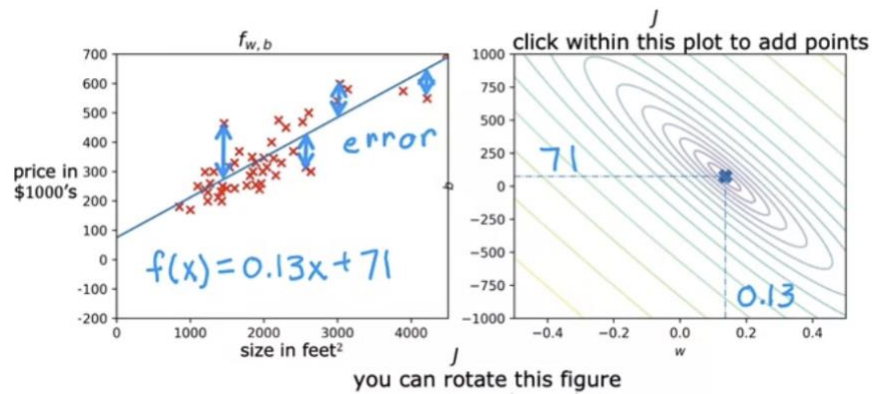
- Poor since the value is far away from the minimum value
- ◆ Another function, $f(x) = 0x + 360$



- ◆ And another function where $f(x) = -0.15x + 500$



- ◆ And another where $f(x) = 0.13x + 71$



- Pretty good fit on contour graph with low error rate in regression plot

Practice Quiz: Regression Model

1.

1 point

For linear regression, the model is $f_{w,b}(x) = wx + b$.

Which of the following are the inputs, or features, that are fed into the model and with which the model is expected to make a prediction?

☐ m

☐ (x, y)

☒ x

☐ w and b .

2. For linear regression, if you find parameters w and b so that $J(w, b)$ is very close to zero, what can you conclude?

1 point

☒ The selected values of the parameters w and b cause the algorithm to fit the training set really well.

☐ The selected values of the parameters w and b cause the algorithm to fit the training set really poorly.

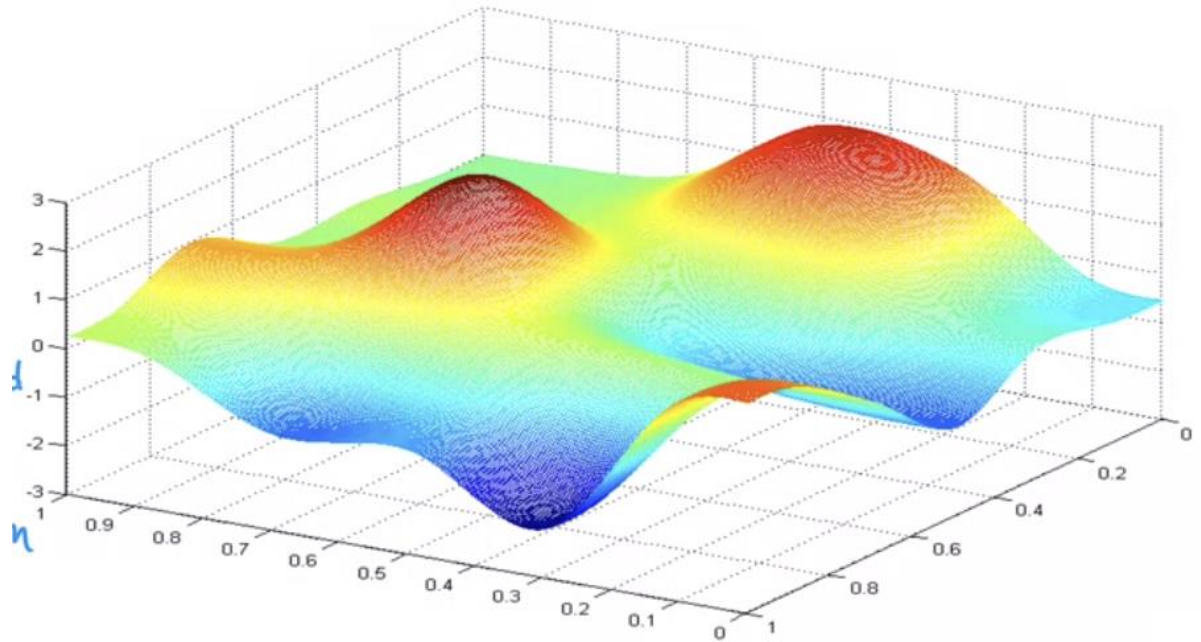
☐ This is never possible -- there must be a bug in the code.

Train the Model with Gradient Descent

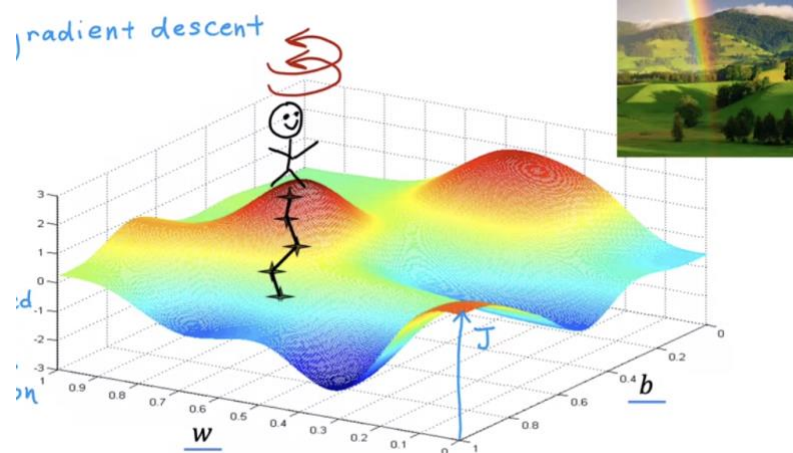
→ Gradient Descent

- ◆ You have a cost function $J(w, b)$ (which could be used to describe any function and not just linear regression), and the goal is to find the $\min J(w, b)$

- Could also be used for more parameters ($J(w_1, w_2, \dots, w_n, b)$)
- ◆ Outline:
 - Start with some value for w and b
 - → Keep changing w, b to reduce $J(w, b)$ until we settle at or near a minimum value
 - ◆ Please note for non-parabolic cost function graphs, there can be more than one minimum value
- ◆ Example:



- Not a squared error cost function and not a linear regression (in linear regression you always end up with a bowl shape)
 - Analogy:

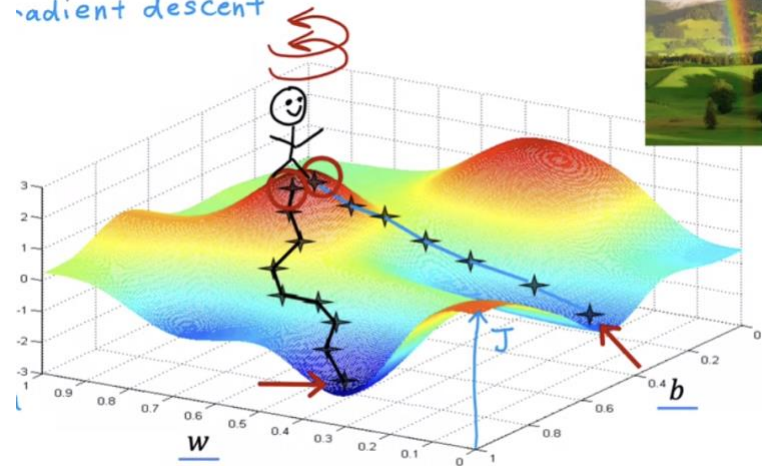


- ◆ Standing on top of a hill and your goal is to get down the hill to the minimum point as fast as possible → you make a 360° turn

and evaluate which is the best way and you walk down a little bit → reevaluate your path again → walk down a little bit, etc.

- This is the same thing the gradient descent algorithm is doing until you reach a local minimum
 - Another path can lead to another local minimum if the starting position is changed

gradient descent



→ Implementing Gradient Descent

- ◆ In gradient descent, the value of w is updated at every step
 - Formula: $w = w \text{ (old value)} - \alpha \times \frac{d}{dw}J(w, b)$
 - $=$ sign in the equation above is the **assignment operator**
 - ◆ This is used in code
 - ◆ Conversely in truth assertion for math you can say $a = c$ to mean a and c are of equal values, but you cannot say things like $a = a + 1$ since that does not make sense in mathematical terms
 - Written as “ $==$ ” in code commonly
 - α is known as the learning rate
 - ◆ Usually between 0 to 1
 - ◆ Denotes how big of a step you take downhill to try and get to the minimum
 - The bigger the alpha value the bigger the steps you are going to be taking
 - $\frac{d}{dw}J(w, b)$ in simple terms is what direction you want to take your step
 - ◆ In combination with α it tells us how big and the direction of the step that we want to take downhill
- ◆ The b value is also updated at every step
 - Formula $b = b \text{ (old value)} - \alpha \times \frac{d}{db}J(w, b)$
- ◆ → The goal is to repeat the steps for the new b value and new w value till convergence is reached
 - Reach the point at a local point where w and b don't change much

- Both w and b are updated **simultaneously**

- Formulas

- ◆ $tmp\ w = w - \alpha \frac{\partial}{\partial w} J(w, b)$

- ◆ $tmp\ b = b - \alpha \frac{\partial}{\partial b} J(w, b)$

- ◆ Note the pre-derivative w goes into the function formula

$$tmp_w = w - \alpha \frac{\partial}{\partial w} J(w, b)$$

$$tmp_b = b - \alpha \frac{\partial}{\partial b} J(w, b)$$

- → Calculate both the $tmp\ w$ and $tmp\ b$ simultaneously and store the values → copy the value of $tmp\ w$ into w ($w = tmp\ w$) and the value of $tmp\ b$ into b ($b = tmp\ b$)

→ Gradient Descent Intuition

- ◆ Goal: repeat until convergence {

$$w = w - \alpha \frac{\partial}{\partial w} J(w, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w, b)$$

- $\alpha = learning\ rate$

- $\frac{\partial}{\partial w} J(w, b) = derivative$

- What is the derivative?

- ◆ You can draw a tangent line at a point on a line and touches the curve

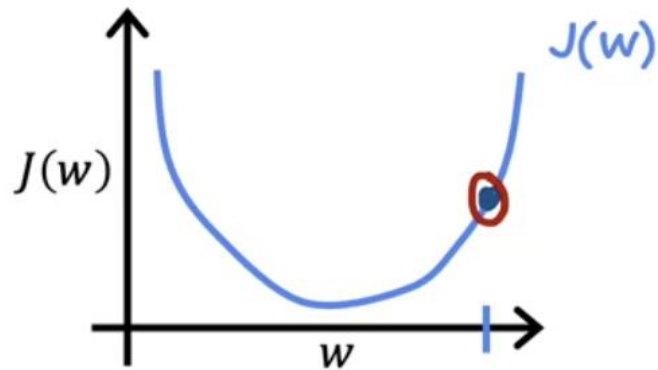
- Example: $J(w)$ cost function

- Gradient descent formula then would be $w = w - \alpha \frac{\partial}{\partial w} J(w)$

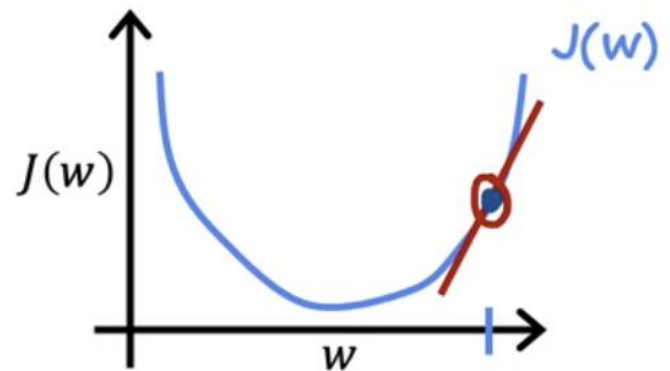
- ◆ Goal is to adjust w to get $\min J(w)$

- Could look at two-dimensional graphs since b can be set to 0

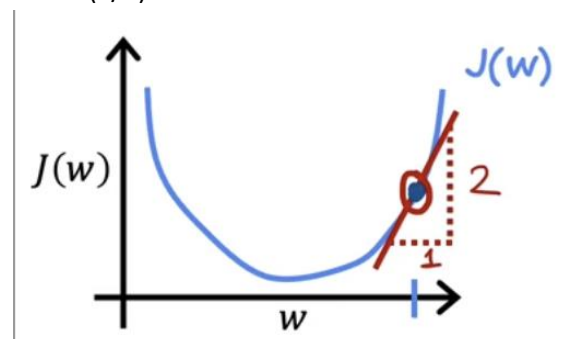
- ◆ → Pick an initial point



- ◆ → update using $w = w - \alpha \frac{\partial}{\partial w} J(w)$
 - Derivative meaning ($\frac{\partial}{\partial w} J(w)$)
 - The slope of the line at a specific point in the derivative of the line



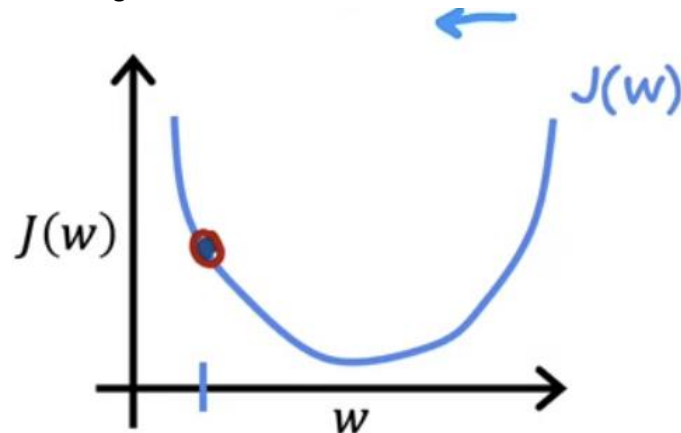
- ◆ To compute the slope, we can use a triangle and divide the height by the width (2/1)



- ◆ In this case the slope is also positive due to its direction, so the derivative is a positive number → $w = w - \alpha \cdot (\text{positive number})$

- Learning rate is always positive
 $\rightarrow w = w - (+) (+) = w - (+)$
 $(+) \rightarrow$ new w value will always be smaller (moving to the left in the case of our example). This makes sense because moving left moves us closer to a minimum $J(w)$ value

- Another example with the same function ($J(w)$ cost function)
 - \rightarrow Starting with a different initial value for w



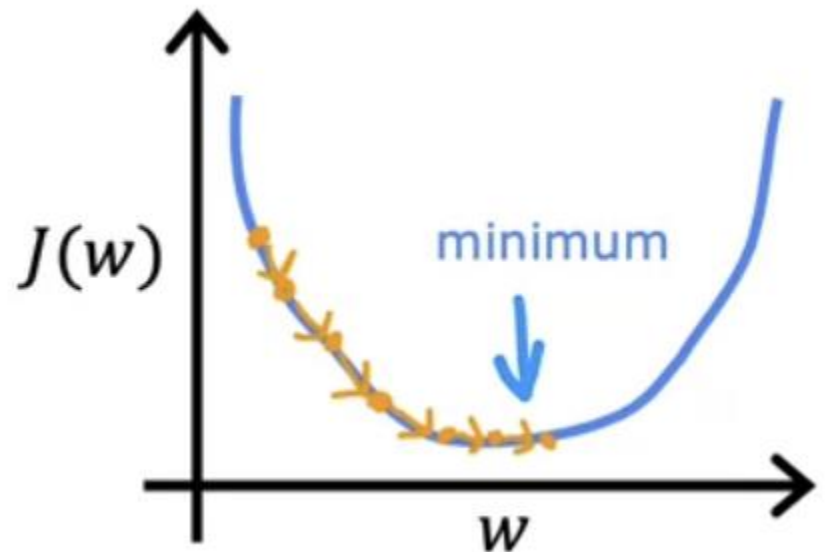
- ◆ Tangent line is negative in this case because it's sloping down and to the right \rightarrow derivative function is negative

- $w = w - \alpha \cdot (\text{negative number})$
 - Learning rate is always positive $\rightarrow w = w - (+) (-) = w - (-) = w + (+)$
 - ◆ New w value in this case will be bigger (moving to the right on the graph). This makes sense because moving right moves us closer to a minimum $J(w)$ value

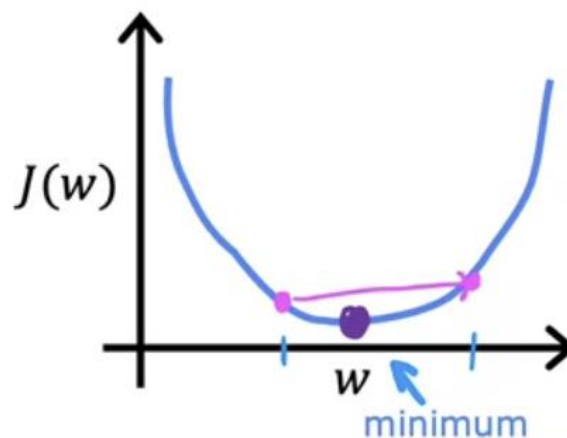
◆ Learning Rate

- Going back to the formula: $w = w - \alpha \frac{\partial}{\partial w} J(w)$
 - α is the learning rate

- ◆ If α is too small, you multiply the derivative term by a small number \rightarrow take a really small step toward the minima

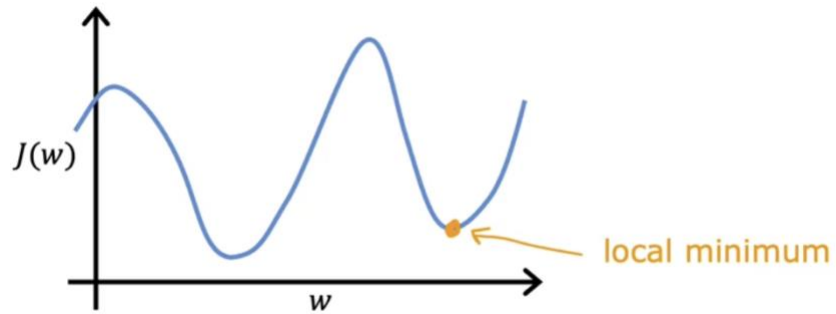


- You eventually get to the minimum but in a lot of steps and really slowly
- ◆ If α is too big, you multiply the derivative term by a big number \rightarrow could potentially overshoot the new w value corresponding to the minima that you get

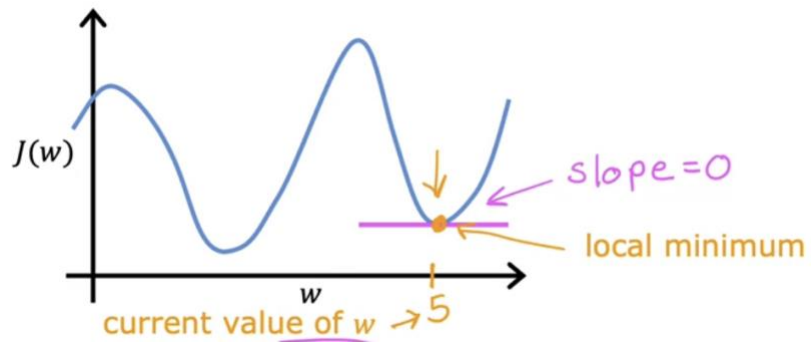


- Cost has actually increased in this example which is not what you want \rightarrow could never reach the minimum and fail to converge or even diverge

- Another example: w is already at a local minimum

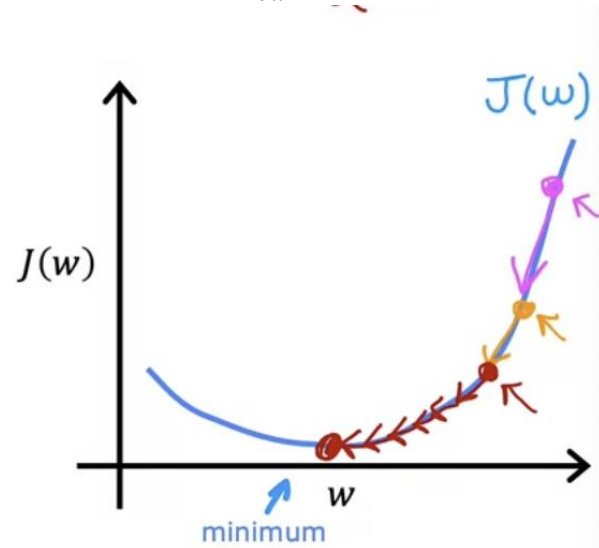


- Not a square error cost function with two local minima
- If you draw a tangent line to the local minima, the slope will equal 0 and the derivative function in $w = w - a \frac{\partial}{\partial w} J(w)$ will also equal 0



- ◆ $\rightarrow w = w - a \cdot 0 \rightarrow w = w$
 - If you are at local minima, gradient descent will leave w unchanged
- **If you reach a local minimum, gradient descent will not work**
 - ◆ Gradient descent can reach a local minimum with a fixed learning rate

- Example: $w = w - \alpha \frac{\partial}{\partial w} J(w)$



- If you start with the first value, there is bigger slope tangent to the point, derivative function will be large
- → Going to the next value for w , the derivative will be smaller, and then smaller as we approach the minimum value, etc.

→ Gradient Descent for Linear Regression

◆ Linear Regression Model

- Function: $f_{w,b}(x) = wx + b$
- Cost Function: $J(w, b) = \frac{1}{2m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})^2$

◆ Gradient Descent Algorithm

- repeat until convergence {

$$w = w - \alpha \frac{\partial}{\partial w} J(w, b) \rightarrow \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)}) x^{(i)} \quad (\text{derivative of cost function with respect to } w)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w, b) \rightarrow \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)}) \quad (\text{derivative of cost function with respect to } b)$$

}

- Could implement gradient descent with the summation formulas

◆ Gradient Descent Algorithm for Linear Regression

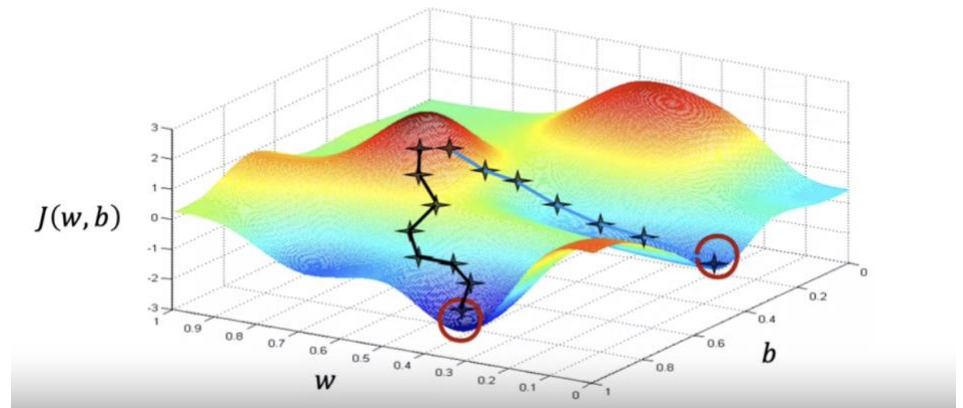
- repeat until convergence {

$$w = w - \alpha \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)}) x^{(i)}$$

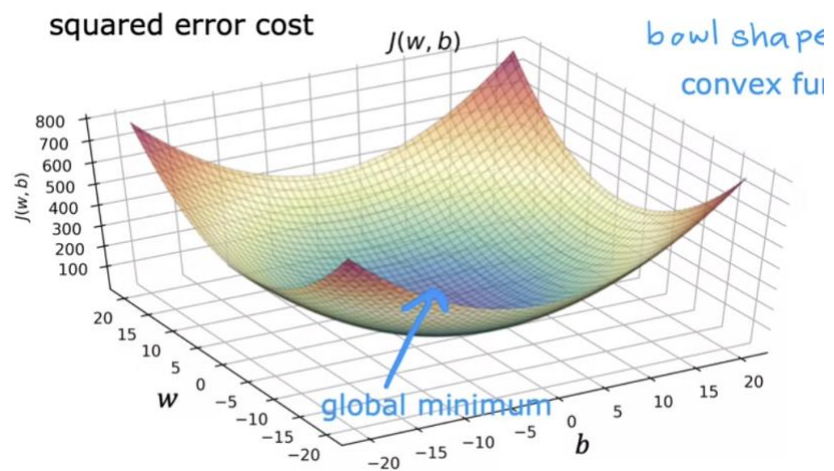
$$b = b - \alpha \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})$$

}

- Want to update w and b simultaneously on each step
- ◆ Gradient Descent Problems
 - Could lead to local minimum rather than global minimum
 - Where you initialize the parameters w and b could lead to different local minima



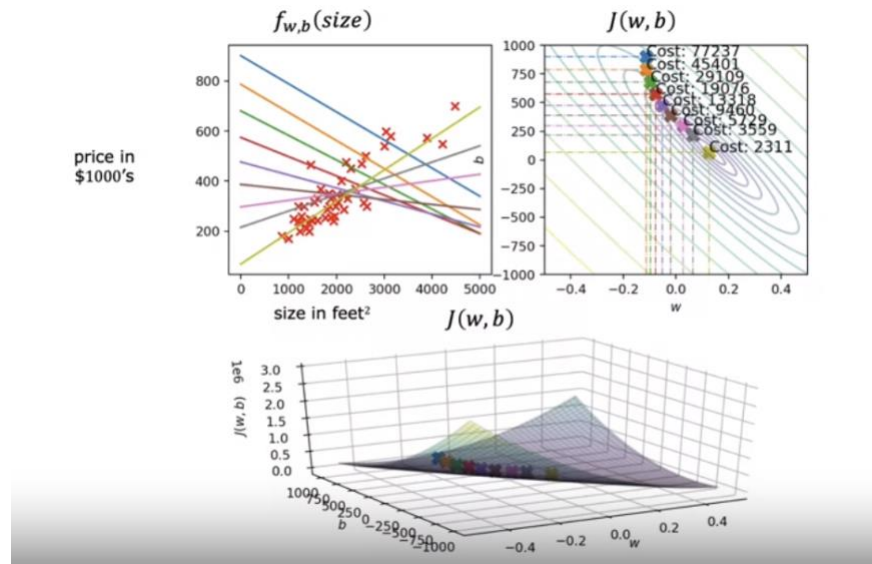
- When using a squared error cost function with linear regression this is not a problem as there is only one global minima due to it being a convex function



→ Running Gradient Descent

- ◆ Example: $f(x) = -0.1x + 900$ for initial starting for a data set

- Cost function slowly moves to decrease the cost as gradient descent continues



- This gradient descent process is called **batch gradient descent**
 - ◆ This means that on every step of gradient descent we are using all the training examples rather than a subset of the training data

→ Practice Quiz: Train the Model with Gradient Descent

1.

1 point

Gradient descent is an algorithm for finding values of parameters w and b that minimize the cost function J .

repeat until convergence {

$$w = w - \alpha \frac{\partial}{\partial w} J(w, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w, b)$$

When $\frac{\partial J(w, b)}{\partial w}$ is a negative number (less than zero), what happens to w after one update step?

- ☒ w increases.
- ☐ w stays the same
- ☐ w decreases
- ☐ It is not possible to tell if w will increase or decrease.

2.

1 point

For linear regression, what is the update step for parameter b ?

- ☒ $b = b - \alpha \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})$
- ☐ $b = b - \alpha \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)}) x^{(i)}$

